

Legal disclosure requirements for an AI wellbeing-call product (US & UK)

Bottom line: a compliant version of this product is realistic to build in both the US and the UK — but only if it discloses that the caller is AI and obtains consent from the elderly recipient (not just the family who pays). No jurisdiction outright bans a benign, human-indistinguishable AI voice, but the entire legal architecture in the US and the transparency regimes in the UK/EU are built around the principle that the person on the line must be able to know they are talking to a machine. The design goal of being "indistinguishable from a human" is precisely what converts a low-risk product into a legally exposed one. The leading competitors that have survived press scrutiny (inTouch, ElderVoice) already disclose AI identity and gate calls on the senior's own consent — which is both the ethical and the legal safe harbor.

The single biggest legal constraint is not a disclosure statute but the **TCPA consent rule**: because a daily AI call uses an "artificial or prerecorded voice," it requires the prior express consent of the person actually being called — and courts have held that a family member cannot supply that consent unless they are the phone's subscriber/customary user or an authorized agent of the recipient. That directly complicates the "family member sets up the service" model.

US federal law: consent is mandatory now; an explicit "this is AI" disclosure is proposed but not yet law

AI-generated voices are already regulated under the TCPA. On February 8, 2024, the FCC released a Declaratory Ruling (CG Docket 23-362, FCC 24-17) confirming that the TCPA's restrictions on "artificial or prerecorded voice" **encompass current AI technologies that generate human voices**, including voice cloning, and that such calls "require the prior express consent of the called party to initiate such calls absent an emergency purpose or exemption." [fcc](#) The ruling was effective on release and explicitly refused to carve out technologies that "purport to provide the equivalent of a live agent." [fcc](#) A conversational AI check-in call is squarely within this definition.

What the TCPA requires today for such calls: (1) prior express consent of the called party; (2) at the **beginning of the message**, a clear statement of the identity of the entity responsible for the call; and (3) during or after the message, the entity's telephone number or address. [cornell](#) [cornell](#) So a caller-identity disclosure ("This is [service name] calling on behalf of your family") is already legally required for artificial-voice calls — but a statement that the voice is AI is not yet.

An explicit AI disclosure is only proposed, not in force. In an NPRM published September 10, 2024 (FCC 24-84), the FCC proposed to require callers to "clearly disclose... at the beginning of each call... that the call is using AI-generated technology," plus AI-specific consent disclosures. [federalregister](#) As of July 2026 this remains a **proposed rule** — the current CFR text (47 CFR § 64.1200) still requires only caller identity and number, not an AI disclosure, and no final Report and Order adopting the AI-disclosure requirement was located. [cornell](#) [federalregister](#) **Ethan should design for this rule as if it will pass**, since the direction of federal policy is unambiguous and the product would be an obvious target.

The telemarketing distinction helps but does not exempt the product. The strictest rules (prior express written consent) attach to calls that "include or introduce an advertisement or constitute telemarketing." [cornell](#) A wellbeing check-in is not telemarketing, which lowers the bar. However, the residential-landline exemption for non-commercial artificial-voice calls only covers callers making **no more than three calls in any 30-day period** — a daily call blows past that limit, re-triggering the consent requirement. [cornell](#) So consent is unavoidable regardless of the non-sales framing.

The consent problem the payer model creates. The TCPA requires consent of the "called party." The FCC and courts define this as the phone's **current subscriber or its customary user**, not merely the intended recipient or a third party who set things up. [epic](#) The Eleventh Circuit (Osorio) held that only the subscriber can consent, "either directly or through an authorized agent," and that a housemate "had no authority to consent in her own right." [findlaw](#) The Ninth Circuit (N.L. v. Credit One) held that "consent from the intended recipient... is not sufficient" — the actual called party must consent. [uscourts](#)

Implication: a family member can only validly consent if the elderly person's phone is on the family member's account (making them the subscriber) or the family member is acting as the recipient's authorized agent. The safest design gets the elderly recipient's own consent — which, notably, is exactly what inTouch does in practice (see below).

US state law: California is the binding constraint; Texas and Florida are not

State coverage is uneven, and — importantly — **several of the newer AI-disclosure statutes are tied to "trade/commerce," "consumer transactions," or "regulated occupations,"** meaning their applicability to a paid-but-non-sales wellbeing call is arguable. The clearest exposure is California.

Law	What it requires	Applicability to this product
CA AB 2905 (eff. ~Jan 1, 2025)	For calls via an automatic dialing-announcing device, a live "unrecorded, natural voice" must state call nature + business identity, ask consent to hear the message, and inform the person if the message uses an "artificial voice" (AI-generated/altered) leginfo getaira	Likely applies; ambiguity is whether a real-time conversational AI counts as a "prerecorded message" under the autodialer statute
Maine 10 M.R.S. § 1500-DD	Prohibits using an AI chatbot in trade/commerce in a way that could mislead a reasonable consumer into thinking they're talking to a human, unless clearly notified they are not; definition explicitly includes "aural" (voice) communication maine	Directly on point for voice AI; turns on whether the call is "trade or commerce with a consumer"
Colorado C.R.S. § 6-1-1704 (eff. June 30, 2026)	Any AI system "intended to interact with consumers" must disclose that the consumer is interacting with an AI system, unless obvious to a reasonable person justia	Applies once effective if a CO resident is called
Utah SB 226 (eff. May 7, 2025)	Must disclose interaction with generative AI (not a human) if the person asks ; stricter proactive disclosure for "high-risk"/regulated-occupation interactions utah	Applies on a prompted basis; a senior asking "are you a real person?" triggers a duty to answer honestly

Law	What it requires	Applicability to this product
CA SB 1001 (BOT Act)	Bot-disclosure duty — but explicitly limited to "online" internet communications and to intent-to-mislead for a sale or election leginfo	Does not apply to voice phone calls

Texas and Florida impose no applicable AI-voice disclosure requirement. Texas SB 140 (effective September 1, 2025) — which the command flagged — turns out on inspection of the enrolled text to contain **no AI/artificial-voice disclosure requirement at all**; it amends telephone-solicitation definitions and seller-registration rules, and "telephone solicitation" is defined as a call to induce a person to "purchase, rent, claim, or receive an item." [texas](#) A non-sales wellbeing call falls outside Chapter 302 entirely. [texas](#) Texas's synthesized-voice consent rules are likewise tied to solicitation contexts. Florida's telemarketing statute prohibits disguising a caller's voice "to defraud, confuse, or injure," and regulates automated sales calls, but has no general AI-disclosure mandate for a benign check-in. [flsenate](#) **So the two named states are not the constraint — California is.**

The California AB 2602 / SB 1381 lead was a false trail. AB 2602 governs contract enforceability for "digital replicas" of performers' voices in entertainment/creative work — not triggered by an AI wellbeing call. [leginfo](#) SB 1381 expands child-sexual-abuse-material criminal law to AI-generated content — wholly irrelevant here. [leginfo](#) The canvas note should be corrected: the relevant California statute is **AB 2905**, not those two.

FTC exposure is about deception, not disclosure per se. The FTC's Impersonation Rule (16 CFR Part 461, effective April 1, 2024) targets falsely posing as a government entity or business [ecfr](#); the Telemarketing Sales Rule governs sales calls. [ecfr](#) The FTC has flagged AI voice cloning as a deception/enforcement priority under its Section 5 authority. [ftc](#) None of this bans a disclosed AI voice, but it confirms that

an undisclosed, human-mimicking AI voice is the fact pattern regulators are hunting.

UK law: no AI-voice-specific rule, but transparency principles apply and PECR probably doesn't bite

The UK has **no Ofcom rule specifically requiring disclosure that a call is AI-generated**; Ofcom's telephone guidance addresses nuisance/silent calls and marketing identity/opt-out, not AI voices. [ofcom](#) Notably, the ICO in September 2025 fined two energy firms for unlawful marketing calls that used "**avatar software**" to make recipients think they were speaking to named humans — a signal that human-mimicking AI voices draw regulatory scrutiny, though that case was a marketing PECR breach. [ico](#)

PECR's automated-call rules likely do not apply, because they are limited to direct marketing. The ICO states that "routine customer service messages do not count as direct marketing" and that the definition doesn't cover live calls for "administrative or customer service purposes." [ico](#) PECR Regulation 19's consent requirement applies to "recorded matter for direct marketing purposes" via automated calling systems. [legislation](#) A genuine, non-promotional wellbeing call therefore falls outside the strict automated-marketing consent regime — a meaningful advantage over the US position. This holds only if the calls never carry promotional content or harvest data for future marketing.

UK GDPR / ICO transparency still applies because the product processes personal (and health-adjacent) data. The ICO requires processing to be "lawful, fair and transparent," and states processing is "unlikely to be considered transparent" if the organisation is not open about how AI is used. [ico](#) [ico](#) Recording conversations requires telling people it's happening. [ico](#) This is a transparency/consent obligation, not an "announce AI" mandate, but disclosing AI identity is the natural way to satisfy fairness for a cognitively vulnerable data subject.

Consumer protection and the Online Safety Act. Under the Digital Markets, Competition and Consumers Act 2024, a "misleading omission" requires an omission of material information within a "commercial practice" affecting a transactional decision [legislation](#) — no authority treats an undisclosed AI voice in a non-marketing call as per se unfair, though aggressive non-disclosure to a vulnerable person could theoretically be argued. The **Online Safety Act 2023 does not apply**: it regulates user-to-user and search internet services and expressly exempts services whose only user content is "one-to-one live aural communications." [legislation](#) [legislation](#)

The "indistinguishable AI" question: no jurisdiction bans it, but disclosure regimes make hiding it the offence

Across the US, UK, and EU, **there is no law making it illegal to use a human-indistinguishable AI voice for a benign purpose per se.** The illegality in every regime attaches to non-disclosure or intent to deceive, not to the quality of the voice:

- US state laws (Maine, Colorado, Utah, CA SB 1001) criminalize/penalize the failure to disclose or intent to mislead — not indistinguishability itself. [maine justia leginfo](#)
- The UK Fraud Act 2006 requires dishonesty plus intent to make a gain or cause loss — a benign check-in lacks those elements. [legislation](#)
- The **EU AI Act Article 50** (transparency, applying from August 2, 2026) requires that AI systems interacting directly with people be designed so persons are "informed that they are interacting with an AI system unless this is obvious." [europa europa](#) An indistinguishable voice is exactly the case where disclosure becomes mandatory rather than optional (because it is no longer "obvious").

The pattern is consistent: the more indistinguishable the voice, the stronger the legal duty to disclose it. Designing for indistinguishability is legally fine; designing to conceal AI identity is what creates liability. For a

cognitively impaired recipient (Ethan's originating use case), non-disclosure also raises the sharpest ethical and consent risk — the 404 Media critique's core concern.

How competitors handle it: the market leaders already disclose

Vendor practice validates the compliant path and shows disclosure is not a market disadvantage:

- **inTouch** discloses explicitly. Its founder states "in the first call... we clearly say this is an AI" and that they disclose "constantly" that "Mary is an AI," and its FAQ states "No calls to seniors ever start without receiving the senior's explicit consent first." [mashable](#) [intouch](#) This consent-gating maps directly onto the TCPA "called-party consent" requirement.
- **ElderVoice** discloses in its FAQ ("Is a real person calling? No, ElderVoice is an advanced Artificial Intelligence") and introduces itself as a companion service "set up by you (the family member)" on the first call. [eldervoice](#)
- **Verocall** states "We are not trying to fool anyone into thinking our agents are human" and even tells families to reconsider if the loved one cannot tell it's not human. [verocall](#) [verocall](#)
- **Joy Calls** identifies itself as AI on its site but appears to rely on family pre-briefing ("talk to your loved one about Joy in advance") rather than a forced start-of-call label, while designing the voice to "sound and feel human." [joycalls](#)
- **CareCall and Senior Protection AI** do not surface any explicit senior-facing AI disclosure or start-of-call script in their public materials — the weakest-disclosure end of the market. [intouch](#)

The 404 Media test underscores why disclosure matters: the reporter's mother perceived the call as "an automated company cold call," and a Harvard ethicist warned AI caregiving must be approached with "humility, honesty." [theverge ispr](#) Caregiver-community sentiment is

actively hostile to undisclosed AI ("I don't want my elderly parents talking to AI"). [reddit](#) Disclosure is thus both a legal hedge and a trust/reputation necessity.

What this means for the build

A compliant product is realistic in both the US and UK, and the compliant design is already the market-leading design.

Concretely, the safe configuration is:

1. **Disclose AI identity at the start of every call** ("Hi, this is [name], an AI assistant calling to check in on behalf of your family"). This is not yet federally mandated but is (a) proposed by the FCC [federalregister](#), (b) required by Maine for voice AI in commerce [maine](#), Colorado from mid-2026 [justia](#), and California's AB 2905 announcement structure [leginfo](#), (c) required under the EU AI Act from August 2026 [europa](#), and (d) demanded by Utah on request [utah](#). One disclosure line satisfies all of them.
2. **Also state caller identity and a callback number** — already required federally for artificial-voice calls today. [cornell](#)
3. **Obtain the elderly recipient's own consent**, not just the family's, because the "called party" is the recipient and third-party consent is legally shaky unless the family member is the phone's subscriber or authorized agent. [findlaw](#) [uscourts](#) This is the deepest tension with the "family buys it" model and needs a real consent-capture step involving the senior.
4. **Keep call content strictly non-promotional** to stay outside US telemarketing rules and UK PECR marketing consent. [ico](#)
5. **Handle recording transparency and health-data processing** under UK GDPR/US state privacy rules (disclose recording; lawful basis for health-adjacent data). [ico](#)

The "indistinguishable from a human" design ambition is the one element to walk back: build a natural, non-robotic voice, but front it with

an honest disclosure and recipient consent. That is legal, ethical, and — per inTouch and ElderVoice — commercially normal.

Where further research would most change the conclusion: First, the **recipient-consent mechanics** deserve a dedicated legal review with a TCPA attorney — specifically whether a workable consent flow can be built when the buyer and the called party differ, and whether the senior's verbal consent captured on the first AI call itself suffices (a chicken-and-egg problem the case law doesn't cleanly resolve). Second, the **FCC AI-disclosure NPRM status** should be re-checked close to launch: if it is finalized, start-of-call AI disclosure shifts from "best practice" to a hard federal requirement, and the exact mandated wording/timing would need to be matched precisely.